

Attention-Fused Multi-Scale OSM Context for Spatio-Temporal GNNs in Urban Mobility Flow Forecasting

Frankline Misango Oyolo
The University of Hong Kong
Hong Kong SAR, China
fmo2704@connect.hku.hk

Julien Coquet
ETH Zürich
Zürich, Switzerland
j.coquet@ethz.ch

Jeffrey Huang
École Polytechnique Fédérale de
Lausanne (EPFL)
Lausanne, Switzerland
jeffrey.huang@epfl.ch

Abstract

Urban mobility systems require accurate short-term forecasts for effective resource management and infrastructure planning. While Graph Neural Networks (GNNs) have advanced spatio-temporal prediction in transportation, they largely neglect explicit multi-scale urban context that drives mobility patterns. We introduce an enhanced Diffusion Convolutional Recurrent Neural Network (DCRNN) that enriches standard architectures with multi-scale OpenStreetMap features and attention-based fusion mechanisms. Our approach integrates comprehensive urban infrastructure characteristics across 500m, 1000m, and 1500m radii, enabling adaptive spatial scale selection based on station characteristics. Evaluated on 90,000 trips across 714 Swiss bike-sharing stations, our method achieves significant improvements over strong baselines (ConvLSTM, ST-GCN): approximately 13% RMSE reduction with RMSE of 2.52, MAE of 1.74, and R^2 of 0.723. The attention mechanism reveals interpretable patterns aligned with urban planning intuitions, where downtown stations emphasize local accessibility while peripheral stations focus on broader connectivity.

CCS Concepts

• **Computing methodologies** → **Neural networks**; *Machine learning*; • **Information systems** → *Geographic information systems*; • **Networks** → Network algorithms.

Keywords

Graph neural networks, spatio-temporal forecasting, urban mobility, attention mechanisms, bike-sharing systems, OpenStreetMap, multi-scale features

1 Introduction

Urban bike-sharing systems generate massive mobility flows that cities must predict to optimize rebalancing operations, plan infrastructure investments, and enhance user experience. The fundamental challenge lies in modeling both spatial dependencies between stations and temporal dynamics simultaneously, while incorporating the rich urban context that fundamentally influences mobility patterns [4].

Current spatio-temporal forecasting methods face a critical limitation: they largely ignore the explicit multi-scale urban context available through comprehensive geographical databases. While Graph Neural Networks have advanced spatial modeling capabilities and Recurrent Neural Networks excel at temporal prediction, existing approaches fail to systematically leverage the hierarchical

urban infrastructure information that drives mobility patterns at different spatial scales.

Our work addresses this gap by enhancing the Diffusion Convolutional Recurrent Neural Network (DCRNN) with two key innovations: (1) systematic multi-scale OpenStreetMap feature extraction capturing urban context at 500m, 1000m, and 1500m radii, and (2) a soft attention fusion mechanism that learns to select appropriate spatial scales for each station based on local urban characteristics.

This approach enables adaptive spatial reasoning where downtown stations can focus on immediate walkability features (500m radius) while peripheral stations emphasize broader connectivity patterns (1500m radius). The result is a more contextually-aware model that better captures the diverse urban morphologies present in modern cities.

1.1 Related Work

Recent advances in spatio-temporal forecasting have explored various neural architectures for urban mobility prediction. ConvLSTM [5] combines convolutional and recurrent mechanisms but treats spatial relationships as Euclidean, missing the complex connectivity patterns of urban networks. ST-GCN [7] applies graph convolution to spatio-temporal data but lacks hierarchical urban context integration.

DCRNN [4] models spatial dependencies through graph diffusion convolution with recurrent units, showing promising results for traffic forecasting. However, existing work focuses primarily on transportation networks without systematic integration of urban geographic features. Recent attention-based approaches [3, 6] have shown benefits of adaptive spatial-temporal modeling, but have not explored multi-scale urban context fusion.

Our approach uniquely combines DCRNN with systematic multi-scale urban features and attention-based fusion specifically designed for bike-sharing prediction, creating a framework that is both theoretically grounded and practically applicable to urban mobility management.

1.2 Contributions

Our main contributions are:

- **Multi-scale urban context integration:** First systematic incorporation of hierarchical OpenStreetMap features at multiple spatial scales (500m/1000m/1500m) into spatio-temporal GNNs.
- **Attention-based scale fusion:** Novel soft attention mechanism that adaptively weights spatial scales based on station urban characteristics and geographic context.

- **Strong empirical validation:** Comprehensive evaluation showing approximately 13% RMSE improvement over strong baselines on 90,000 trips across 714 stations.
- **Interpretable urban insights:** Analysis demonstrating that learned attention patterns align with urban planning principles and provide actionable insights for infrastructure development.

2 Multi-Scale Urban Context Framework

2.1 Problem Formulation

Given a bike-sharing network represented as a graph $G = (V, E, A)$ with $N = 714$ stations, our objective is to predict the future flow matrix $\hat{Y}^{(t+1)} \in \mathbb{R}^{N \times N}$ from historical sequences $H^{(t-T:t)}$ and multi-scale urban features X .

The spatial adjacency matrix A encodes station proximity using Gaussian kernel weighting:

$$A_{ij} = \exp\left(-\frac{d_{ij}^2}{\sigma^2}\right) \cdot \mathbf{1}_{j \in \mathcal{N}_k(i)} \quad (1)$$

where d_{ij} is the haversine distance between stations i and j , $\sigma = 2$ km is the bandwidth parameter, and $\mathcal{N}_k(i)$ represents the $k = 10$ nearest neighbors of station i .

Each station v_i is enriched with hierarchical OpenStreetMap-derived features extracted at multiple spatial scales:

$$X_i^{(r)} = \text{OSMExtract}(\text{lat}_i, \text{lon}_i, \text{radius} = r) \quad (2)$$

$$= [f_{\text{transport}}^{(r)}, f_{\text{poi}}^{(r)}, f_{\text{landuse}}^{(r)}, f_{\text{infra}}^{(r)}] \quad (3)$$

For $r \in \{500\text{m}, 1000\text{m}, 1500\text{m}\}$, this yields 47 features per scale encompassing transportation infrastructure (road density, intersection count), points-of-interest diversity (restaurants, retail, education), land use patterns (residential, commercial ratios), and network connectivity metrics.

2.2 Attention-Based Multi-Scale Feature Fusion

Rather than simple concatenation of multi-scale features, we employ learnable attention with station-specific context awareness. The fusion mechanism computes attention weights based on both the features themselves and the station's urban context:

$$c_i = \tanh(W_c \cdot [\text{lat}_i, \text{lon}_i, \text{degree}_i]) \quad (4)$$

$$\alpha_i^{(r)} = \text{softmax}(W_a \cdot [X_i^{(500)}, X_i^{(1000)}, X_i^{(1500)}, c_i]) \quad (5)$$

$$X_i^{\text{fused}} = \sum_r \alpha_i^{(r)} \cdot X_i^{(r)} \quad (6)$$

Here, c_i captures station centrality characteristics and degree_i represents the node degree in the spatial graph. This design enables urban core stations (high centrality) to emphasize fine-grained walkability features (500m) while peripheral stations (low centrality) focus on broader accessibility patterns (1500m).

2.3 Enhanced DCRNN Architecture

We integrate the fused multi-scale features into DCRNN's Diffusion Convolutional GRU cells. The diffusion convolution captures spatial

dependencies through bidirectional random walks on the station graph:

$$\mathcal{D}_G(X, \Theta) = \sum_{k=0}^{K-1} \Theta_k^{(\text{out})} \mathbf{P}_k^{(\text{out})} X + \Theta_k^{(\text{in})} \mathbf{P}_k^{(\text{in})} X \quad (7)$$

$$\mathbf{r}^{(t)} = \sigma(\mathcal{D}_G([H^{(t-1)}, X^{\text{fused}}], \Theta_r)) \quad (8)$$

$$\mathbf{u}^{(t)} = \sigma(\mathcal{D}_G([H^{(t-1)}, X^{\text{fused}}], \Theta_u)) \quad (9)$$

$$\tilde{H}^{(t)} = \tanh(\mathcal{D}_G([X^{(t)}, \mathbf{r}^{(t)} \odot H^{(t-1)}], \Theta_h)) \quad (10)$$

$$H^{(t)} = \mathbf{u}^{(t)} \odot H^{(t-1)} + (1 - \mathbf{u}^{(t)}) \odot \tilde{H}^{(t)} \quad (11)$$

where $\mathbf{P}_k^{(\text{out})} = (\mathbf{D}^{-1}\mathbf{A})^k$ and $\mathbf{P}_k^{(\text{in})} = (\mathbf{D}^{-1}\mathbf{A}^T)^k$ capture forward and backward diffusion processes with $K = 3$ diffusion steps. The encoder processes 6-hour historical sequences while the decoder predicts 1-hour ahead flows using scheduled sampling with probability $\epsilon_t = \max(0, 1 - t/T)$.

3 Experimental Setup and Results

3.1 Dataset and Preprocessing

Our evaluation leverages a comprehensive Swiss bike-sharing dataset with the following characteristics:

- **Temporal Coverage:** 192 hours (8 days) with 98.7% data completeness
- **Spatial Scale:** 714 stations across 3 major Swiss cities with mean inter-station distance of 1.2km
- **Flow Statistics:** 90,000 trips with mean flow of 12.6 and standard deviation of 18.3 trips/hour/station-pair
- **Feature Dimensionality:** 141-dimensional feature vectors per station (47 features $\times$ 3 scales)
- **Network Properties:** Mean degree 10.0, clustering coefficient 0.31, network diameter 8 hops

The OpenStreetMap feature extraction employs the OSMnx framework [1] to systematically capture urban infrastructure across four categories: transportation infrastructure, points-of-interest, land use patterns, and connectivity metrics. The 8-day timespan captures essential mobility variations including weekday/weekend patterns (correlation $\rho = 0.34$), peak hour dynamics (morning: 7-9am, evening: 5-7pm), and weather impacts (rain days show 23% flow reduction).

3.2 Training Configuration and Baseline Methods

Model training employs Adam optimization (learning rate 10^{-3} , $\beta_1 = 0.9$, $\beta_2 = 0.999$), batch size 32, hidden dimensions 64, dropout 0.2, and L2 regularization $\lambda = 10^{-4}$. Early stopping with patience=15 prevents overfitting based on validation RMSE.

We compare against comprehensive baselines representing different methodological approaches:

- **Historical Average:** Simple temporal averaging baseline
- **ARIMA:** Auto-regressive integrated moving average for time series prediction
- **XGBoost:** Gradient boosting with engineered spatiotemporal features [2]

- **ConvLSTM**: Convolutional LSTM for spatio-temporal modeling [5]
- **ST-GCN**: Spatio-temporal graph convolutional networks [7]

These baselines represent traditional statistical, machine learning, and state-of-the-art deep learning approaches to spatio-temporal prediction.

3.3 Performance Comparison

Table 1 presents the comparative performance showing substantial improvements from our multi-scale approach.

Table 1: Performance comparison on Swiss bike-sharing dataset

Method	RMSE	MAE	R ²
Historical Average	4.23	3.12	0.421
ARIMA	3.87	2.89	0.534
XGBoost	3.34	2.43	0.598
ConvLSTM	3.12	2.19	0.634
ST-GCN	2.89	1.97	0.672
Our Method	2.52	1.74	0.723
Improvement vs ST-GCN	12.8%	11.7%	7.6%

The results demonstrate consistent and significant improvements across all metrics. Statistical validation through paired t-tests confirms significance ($p < 0.001$, Cohen’s $d = 1.34$). Performance gains are consistent across different station types: urban core stations (+15% vs. ST-GCN), suburban stations (+12%), and mixed urban zones (+14%).

3.4 Ablation Study

Table 2 reveals the progressive contribution of each component in our framework.

Table 2: Ablation study showing component contributions

Configuration	RMSE	MAE	R ²
DCRNN baseline	3.21	2.34	0.598
+ Single-scale OSM (1000m)	2.78	1.89	0.687
+ Multi-scale (no attention)	2.64	1.81	0.704
+ Attention fusion (full)	2.52	1.74	0.723

The ablation study demonstrates that OpenStreetMap features provide the largest performance gain (13.4% RMSE reduction), multi-scale integration adds 5.0

3.5 Attention Pattern Analysis

Analysis of learned attention weights reveals three distinct station clusters with interpretable urban characteristics:

- **Urban Dense (n=187)**: High attention on 500m features ($\alpha_{500} = 0.68 \pm 0.12$), emphasizing local walkability and immediate amenity access.

- **Peripheral (n=342)**: Strong weighting on 1500m features ($\alpha_{1500} = 0.59 \pm 0.15$), focusing on broader connectivity and accessibility patterns.
- **Mixed Zones (n=185)**: Balanced attention weights across scales, reflecting intermediate urban morphology and transitional characteristics.

Statistical correlation analysis shows attention weights strongly correlate with urban density metrics ($r = 0.74$, $p < 0.001$) and negatively with distance-to-city-center ($r = -0.68$). This validates our hypothesis that urban morphology drives optimal spatial scale selection.

4 Urban Planning Insights and Applications

4.1 Interpretable Spatial Patterns

The learned attention mechanism provides actionable insights for urban planning and infrastructure development. Downtown stations’ emphasis on 500m features aligns with walkability principles and 15-minute city concepts, while peripheral stations’ focus on 1500m connectivity reflects car-dependent suburban mobility patterns.

These patterns suggest targeted infrastructure strategies: dense urban areas benefit from local amenity integration and pedestrian infrastructure, while suburban stations require enhanced connectivity through transit links and cycling infrastructure extending to broader catchment areas.

4.2 Operational Applications

The model’s computational efficiency (0.12s inference time for 714 stations) enables real-time deployment for operational decision-making. Applications include:

- **Dynamic Rebalancing**: Proactive truck deployment based on predicted flow imbalances
- **Pricing Optimization**: Dynamic pricing strategies based on anticipated demand patterns
- **Infrastructure Planning**: Station placement optimization using attention weight analysis
- **Capacity Management**: Proactive capacity adjustments to reduce user wait times

4.3 Cross-City Generalization

Cross-validation across cities demonstrates consistent performance (Zurich: RMSE 2.41, Basel: 2.58, Bern: 2.67), suggesting good transferability despite different urban morphologies and system characteristics. Transfer learning validation shows that pre-training on one city and fine-tuning on another requires only 15% of original training time while achieving 94% of full training performance.

5 Discussion and Future Directions

5.1 Technical Contributions and Limitations

Our framework introduces systematic multi-scale urban context integration and attention-based fusion for spatio-temporal GNNs. The approach maintains computational efficiency while providing substantial performance improvements and interpretable insights aligned with urban planning principles.

Current limitations include reliance on static OpenStreetMap features that miss dynamic urban factors such as weather conditions, public events, and real-time transit disruptions. The 8-day dataset, while comprehensive, represents limited seasonal variation and may not capture long-term mobility pattern evolution.

5.2 Broader Impact for Smart Cities

This work contributes to the development of more intelligent and responsive urban mobility systems. Enhanced prediction accuracy supports sustainable transportation through improved bike-sharing efficiency, potentially encouraging mode shift from private vehicles and reducing urban emissions.

The interpretable attention mechanism provides transparency for urban planners and transportation operators, supporting evidence-based infrastructure decisions and policy development. However, algorithmic bias toward high-demand areas requires careful consideration to ensure equitable service provision across different urban communities.

5.3 Future Research Directions

Future extensions include: (1) integration of dynamic external factors (weather, events, transit disruptions), (2) temporal attention mechanisms for adaptive historical context weighting, (3) graph structure learning for evolving urban networks, (4) multi-modal integration extending to scooters and car-sharing systems, and (5) causal modeling approaches for policy-relevant counterfactual analysis.

6 Conclusion

We present an attention-fused multi-scale framework that significantly advances spatio-temporal GNN performance for urban mobility forecasting. By systematically integrating hierarchical OpenStreetMap features with learnable attention fusion, our approach achieves approximately 13% RMSE improvement over strong baselines while providing interpretable insights aligned with urban planning principles.

The key innovation lies in recognizing that effective urban mobility prediction requires adaptive spatial context—downtown stations benefit from local accessibility features while peripheral stations require broader connectivity information. This insight, validated through comprehensive evaluation on 90,000 trips, offers a promising direction for developing more intelligent and contextually-aware urban mobility systems.

Our results demonstrate that incorporating explicit multi-scale urban context through attention-based feature fusion provides a strong foundation for next-generation smart city applications, supporting both operational efficiency and strategic infrastructure planning in urban transportation systems.

Acknowledgments

This work was conducted within the Blue City Project at EPFL. We thank the Swiss bike-sharing operators for data access and the OpenStreetMap community for maintaining comprehensive urban infrastructure data.

References

- [1] Geoff Boeing. 2017. OSMnx: New methods for acquiring, constructing, analyzing, and visualizing complex street networks. *Computers, Environment and Urban Systems* 65 (2017), 126–139.
- [2] Tianqi Chen and Carlos Guestrin. 2016. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 785–794.
- [3] Shengnan Guo, Youfang Lin, Ning Feng, Chao Song, and Huaiyu Wan. 2020. Attention based spatial-temporal graph convolutional networks for traffic flow forecasting. In *Proceedings of the 33rd AAAI Conference on Artificial Intelligence*. 922–929.
- [4] Yaguang Li, Rose Yu, Cyrus Shahabi, and Yan Liu. 2018. Diffusion convolutional recurrent neural network: Data-driven traffic forecasting. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- [5] Xingjian Shi, Zhourong Chen, Hao Wang, Dit-Yan Yeung, Wai-Kin Wong, and Wang-chun Woo. 2015. Convolutional LSTM network: A machine learning approach for precipitation nowcasting. In *Advances in Neural Information Processing Systems*. 802–810.
- [6] Zonghan Wu, Shirui Pan, Guodong Long, Jing Jiang, and Chengqi Zhang. 2019. Graph wavenet for deep spatial-temporal graph modeling. In *Proceedings of the 28th International Joint Conference on Artificial Intelligence*. 1907–1913.
- [7] Bing Yu, Haoteng Yin, and Zhanxing Zhu. 2018. Spatio-temporal graph convolutional networks: A deep learning framework for traffic forecasting. In *Proceedings of the 27th International Joint Conference on Artificial Intelligence*. 3634–3640.